

practice project

July 12, 2022

```
[1]: #load the file using pandas
      #import all required libraries
      import pandas as pd
      import numpy as np
      import matplotlib.pyplot as plt
      import seaborn as sns
```

```
[2]: data=pd.read_csv('hour.csv')
      data
```

```
[2]:      instant      dteday  season  yr  mnth  hr  holiday  weekday  \
0          1  2011-01-01        1   0    1    0         0         6
1          2  2011-01-01        1   0    1    1         0         6
2          3  2011-01-01        1   0    1    2         0         6
3          4  2011-01-01        1   0    1    3         0         6
4          5  2011-01-01        1   0    1    4         0         6
...      ...      ...      ..   ..   ...   ...   ...
17374    17375  2012-12-31        1   1    12   19         0         1
17375    17376  2012-12-31        1   1    12   20         0         1
17376    17377  2012-12-31        1   1    12   21         0         1
17377    17378  2012-12-31        1   1    12   22         0         1
17378    17379  2012-12-31        1   1    12   23         0         1

      workingday  weathersit  temp  atemp  hum  windspeed  casual  \
0              0          1  0.24  0.2879  0.81    0.0000         3
1              0          1  0.22  0.2727  0.80    0.0000         8
2              0          1  0.22  0.2727  0.80    0.0000         5
3              0          1  0.24  0.2879  0.75    0.0000         3
4              0          1  0.24  0.2879  0.75    0.0000         0
...      ...      ...      ...   ...   ...   ...
17374          1          2  0.26  0.2576  0.60    0.1642        11
17375          1          2  0.26  0.2576  0.60    0.1642         8
17376          1          1  0.26  0.2576  0.60    0.1642         7
17377          1          1  0.26  0.2727  0.56    0.1343        13
17378          1          1  0.26  0.2727  0.65    0.1343        12

      registered  cnt
```

```

0          13   16
1          32   40
2          27   32
3          10   13
4           1    1
...
17374      108  119
17375       81   89
17376       83   90
17377       48   61
17378       37   49

```

[17379 rows x 17 columns]

```
[3]: data.head()
```

```

[3]:   instant    dteday  season  yr  mnth  hr  holiday  weekday  workingday  \
0      1  2011-01-01        1  0    1    0         0         6         0
1      2  2011-01-01        1  0    1    1         0         6         0
2      3  2011-01-01        1  0    1    2         0         6         0
3      4  2011-01-01        1  0    1    3         0         6         0
4      5  2011-01-01        1  0    1    4         0         6         0

   weathersit  temp  atemp  hum  windspeed  casual  registered  cnt
0          1  0.24  0.2879  0.81         0.0         3         13   16
1          1  0.22  0.2727  0.80         0.0         8         32   40
2          1  0.22  0.2727  0.80         0.0         5         27   32
3          1  0.24  0.2879  0.75         0.0         3         10   13
4          1  0.24  0.2879  0.75         0.0         0          1    1

```

```

[5]: #shape
data.shape

```

[5]: (17379, 17)

```

[6]: #null values
data.isnull()

```

```

[6]:   instant  dteday  season  yr  mnth  hr  holiday  weekday  \
0      False  False  False  False  False  False  False  False
1      False  False  False  False  False  False  False  False
2      False  False  False  False  False  False  False  False
3      False  False  False  False  False  False  False  False
4      False  False  False  False  False  False  False  False
...
17374  False  False  False  False  False  False  False  False
17375  False  False  False  False  False  False  False  False

```

17376	False	False	False	False	False	False	False	False
17377	False	False	False	False	False	False	False	False
17378	False	False	False	False	False	False	False	False

	workingday	weathersit	temp	atemp	hum	windspeed	casual	\
0	False	False	False	False	False	False	False	
1	False	False	False	False	False	False	False	
2	False	False	False	False	False	False	False	
3	False	False	False	False	False	False	False	
4	False	False	False	False	False	False	False	
...	
17374	False	False	False	False	False	False	False	
17375	False	False	False	False	False	False	False	
17376	False	False	False	False	False	False	False	
17377	False	False	False	False	False	False	False	
17378	False	False	False	False	False	False	False	

	registered	cnt
0	False	False
1	False	False
2	False	False
3	False	False
4	False	False
...
17374	False	False
17375	False	False
17376	False	False
17377	False	False
17378	False	False

[17379 rows x 17 columns]

```
[7]: data.isnull().sum()
```

```
[7]: instant      0
      dteday      0
      season      0
      yr          0
      mnth        0
      hr          0
      holiday      0
      weekday      0
      workingday   0
      weathersit    0
      temp        0
      atemp       0
      hum         0
```

```
windspeed      0
casual          0
registered      0
cnt             0
dtype: int64
```

```
[8]: #summary
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 17379 entries, 0 to 17378
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   instant     17379 non-null  int64
1   dteday      17379 non-null  object
2   season      17379 non-null  int64
3   yr          17379 non-null  int64
4   mnth        17379 non-null  int64
5   hr          17379 non-null  int64
6   holiday      17379 non-null  int64
7   weekday     17379 non-null  int64
8   workingday  17379 non-null  int64
9   weathersit   17379 non-null  int64
10  temp        17379 non-null  float64
11  atemp       17379 non-null  float64
12  hum         17379 non-null  float64
13  windspeed   17379 non-null  float64
14  casual      17379 non-null  int64
15  registered  17379 non-null  int64
16  cnt         17379 non-null  int64
dtypes: float64(4), int64(12), object(1)
memory usage: 2.3+ MB
```

```
[9]: data.describe()
```

```
[9]:
```

	instant	season	yr	mnth	hr \
count	17379.0000	17379.000000	17379.000000	17379.000000	17379.000000
mean	8690.0000	2.501640	0.502561	6.537775	11.546752
std	5017.0295	1.106918	0.500008	3.438776	6.914405
min	1.0000	1.000000	0.000000	1.000000	0.000000
25%	4345.5000	2.000000	0.000000	4.000000	6.000000
50%	8690.0000	3.000000	1.000000	7.000000	12.000000
75%	13034.5000	3.000000	1.000000	10.000000	18.000000
max	17379.0000	4.000000	1.000000	12.000000	23.000000

	holiday	weekday	workingday	weathersit	temp \
--	---------	---------	------------	------------	--------

count	17379.000000	17379.000000	17379.000000	17379.000000	17379.000000
mean	0.028770	3.003683	0.682721	1.425283	0.496987
std	0.167165	2.005771	0.465431	0.639357	0.192556
min	0.000000	0.000000	0.000000	1.000000	0.020000
25%	0.000000	1.000000	0.000000	1.000000	0.340000
50%	0.000000	3.000000	1.000000	1.000000	0.500000
75%	0.000000	5.000000	1.000000	2.000000	0.660000
max	1.000000	6.000000	1.000000	4.000000	1.000000

	atemp	hum	windspeed	casual	registered \
count	17379.000000	17379.000000	17379.000000	17379.000000	17379.000000
mean	0.475775	0.627229	0.190098	35.676218	153.786869
std	0.171850	0.192930	0.122340	49.305030	151.357286
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.333300	0.480000	0.104500	4.000000	34.000000
50%	0.484800	0.630000	0.194000	17.000000	115.000000
75%	0.621200	0.780000	0.253700	48.000000	220.000000
max	1.000000	1.000000	0.850700	367.000000	886.000000

	cnt
count	17379.000000
mean	189.463088
std	181.387599
min	1.000000
25%	40.000000
50%	142.000000
75%	281.000000
max	977.000000

```
[10]: #check duplicate data
duplicate=data[data.duplicated()]
duplicate
```

```
[10]: Empty DataFrame
Columns: [instant, dteday, season, yr, mnth, hr, holiday, weekday, workingday,
weathersit, temp, atemp, hum, windspeed, casual, registered, cnt]
Index: []
```

```
[11]: #Sanity checks
#Check if registered + casual = cnt for all the records. The two must add to
→cnt,
#if not the row is junk and should be dropped.
data['registered']+data['casual']!=data['cnt']
```

```
[11]: 0      False
      1      False
      2      False
```

```

3      False
4      False
...
17374   False
17375   False
17376   False
17377   False
17378   False
Length: 17379, dtype: bool

```

```
[3]: #Sum
np.sum(data['registered']+data['casual']!=data['cnt'])
```

```
[3]: 0
```

```
[4]: #suppose reg+casual!=cnt, then the code will be
data.drop(data[data['registered']+data['casual']!=data['cnt']],
          ↪index,inplace=True)
```

```
[16]: #Month values should be 1-12 only
data['mnth'].unique()
```

```
[16]: array([ 1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12])
```

```
[17]: #Hour should be 0-23
data['hr'].unique()
```

```
[17]: array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
          17, 18, 19, 20, 21, 22, 23])
```

```
[3]: #Variables 'casual', 'registered' are redundant and need to be dropped.
# 'Instant' is the index, and needs to be dropped too.
# The date column dteday will not be used in the model building, and hence needs
  ↪to be dropped.
# Create new dataframe named 'inp1'.
list=['casual','registered','dteday','instant']
inp1=data.drop(list,axis=1).copy()
```

```
[4]: inp1
```

```
[4]:
```

	season	yr	mnth	hr	holiday	weekday	workingday	weathersit	temp \
0	1	0	1	0	0	6	0	1	0.24
1	1	0	1	1	0	6	0	1	0.22
2	1	0	1	2	0	6	0	1	0.22
3	1	0	1	3	0	6	0	1	0.24
4	1	0	1	4	0	6	0	1	0.24
...

17374	1	1	12	19	0	1	1	2	0.26
17375	1	1	12	20	0	1	1	2	0.26
17376	1	1	12	21	0	1	1	1	0.26
17377	1	1	12	22	0	1	1	1	0.26
17378	1	1	12	23	0	1	1	1	0.26

	atemp	hum	windspeed	cnt
0	0.2879	0.81	0.0000	16
1	0.2727	0.80	0.0000	40
2	0.2727	0.80	0.0000	32
3	0.2879	0.75	0.0000	13
4	0.2879	0.75	0.0000	1
...
17374	0.2576	0.60	0.1642	119
17375	0.2576	0.60	0.1642	89
17376	0.2576	0.60	0.1642	90
17377	0.2727	0.56	0.1343	61
17378	0.2727	0.65	0.1343	49

[17379 rows x 13 columns]

```
[21]: inp1.shape
```

```
[21]: (17379, 13)
```

```
[22]: #Univariate analysis -
#Describe the numerical fields in the dataset using pandas describe method
inp1.describe()
```

```
[22]:
```

	season	yr	mnth	hr	holiday \
count	17379.000000	17379.000000	17379.000000	17379.000000	17379.000000
mean	2.501640	0.502561	6.537775	11.546752	0.028770
std	1.106918	0.500008	3.438776	6.914405	0.167165
min	1.000000	0.000000	1.000000	0.000000	0.000000
25%	2.000000	0.000000	4.000000	6.000000	0.000000
50%	3.000000	1.000000	7.000000	12.000000	0.000000
75%	3.000000	1.000000	10.000000	18.000000	0.000000
max	4.000000	1.000000	12.000000	23.000000	1.000000


	weekday	workingday	weathersit	temp	atemp \
count	17379.000000	17379.000000	17379.000000	17379.000000	17379.000000
mean	3.003683	0.682721	1.425283	0.496987	0.475775
std	2.005771	0.465431	0.639357	0.192556	0.171850
min	0.000000	0.000000	1.000000	0.020000	0.000000
25%	1.000000	0.000000	1.000000	0.340000	0.333300
50%	3.000000	1.000000	1.000000	0.500000	0.484800
75%	5.000000	1.000000	2.000000	0.660000	0.621200

max	6.000000	1.000000	4.000000	1.000000	1.000000
-----	----------	----------	----------	----------	----------


	hum	windspeed	cnt
count	17379.000000	17379.000000	17379.000000
mean	0.627229	0.190098	189.463088
std	0.192930	0.122340	181.387599
min	0.000000	0.000000	1.000000
25%	0.480000	0.104500	40.000000
50%	0.630000	0.194000	142.000000
75%	0.780000	0.253700	281.000000
max	1.000000	0.850700	977.000000

```
[23]: #Make density plot for temp.
      #This would give a sense of the centrality and the spread of the distribution.
      sns.kdeplot(data['temp'])
```


```
[23]: <AxesSubplot:xlabel='temp', ylabel='Density'>
```



```
[29]: #Boxplot for atemp.
      #Are there any outliers?
      sns.boxplot(data=inp1.atemp)
      plt.show()
      #there are no outliers
```





```
[14]: #Histogram for hum
      #Do you detect any abnormally high values?
      inpl.hum.plot.hist(bins=30)#Gives clear value if no.of bins increased
      plt.show()
```




```
[13]: sns.histplot(data=inp1.hum) #better plot is produced using seaborn
```

```
[13]: <AxesSubplot:xlabel='hum', ylabel='Count'>
```




```
[15]: #Density plot for windspeed  
sns.kdeplot(data['windspeed'])
```

```
[15]: <AxesSubplot:xlabel='windspeed', ylabel='Density'>
```



```
[20]: #Box and density plot for cnt - this is the variable of interest.
      #Do you see any outliers in the boxplot?
      #Does the density plot provide a similar insight?
      inpl.cnt.plot.density()
```


```
[20]: <AxesSubplot:ylabel='Density'>
```



```
[21]: sns.boxplot(inp1.cnt)
plt.show()
#we have a lot of outliers present at the higher side
```

/usr/local/lib/python3.7/site-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning



```
[22]: #Outlier treatment -
#Cnt - looks like some hours have rather high values of cnt. We'll need to
→ treat these outliers so that they don't skew our analysis and our model.
#Find out the following percentiles - 10, 25, 50, 75, 90, 95, 99
#Decide the cutoff percentile and drop records with values higher than the
→ cutoff.
#Name the new dataframe 'inp2'.
inp1.cnt.quantile([0.1,0.25,0.5,0.75,0.90,0.95,0.99])
```

```
[22]: 0.10      9.00
      0.25     40.00
      0.50    142.00
```


```
0.75    281.00
0.90    451.20
0.95    563.10
0.99    782.22
Name: cnt, dtype: float64
```

```
[5]: #Decide the cutoff percentile and drop records with values higher than the
      ↪cutoff.
      #Name the new dataframe 'inp2'.
      #The cut off is decided at 95 percentile
      inp2=inp1[inp1.cnt<563].copy()
```

```
[26]: sns.boxplot(inp2.cnt)
      plt.show()
```

/usr/local/lib/python3.7/site-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning




```
[31]: #Bi-variate analysis
      #Make box plot for cnt vs hr
      #What kind of pattern do you see?
      sns.boxplot('hr', 'cnt', data=inp2)
```

```
plt.figure(figsize=[12,6])
```

/usr/local/lib/python3.7/site-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

[31]: <Figure size 864x432 with 0 Axes>




<Figure size 864x432 with 0 Axes>

```
[32]: #Make boxplot for cnt vs weekday
sns.boxplot('weekday', 'cnt', data=inp2)
plt.figure(figsize=[12,6])
```

/usr/local/lib/python3.7/site-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

[32]: <Figure size 864x432 with 0 Axes>




<Figure size 864x432 with 0 Axes>

```
[33]: #Make boxplot for cnt vs month
sns.boxplot('mnth','cnt',data=inp2)
plt.figure(figsize=[12,6])
```

/usr/local/lib/python3.7/site-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

[33]: <Figure size 864x432 with 0 Axes>




<Figure size 864x432 with 0 Axes>

```
[34]: #Make boxplot for cnt vs season
sns.boxplot('season', 'cnt', data=inp2)
plt.figure(figsize=[12,6])
```

/usr/local/lib/python3.7/site-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

[34]: <Figure size 864x432 with 0 Axes>




<Figure size 864x432 with 0 Axes>

```
[35]: #Make a bar plot with the median value of cnt for each hr
      #Does this paint a different picture than the box plot?
      sns.barplot('hr', 'cnt', data=inp2)
```

/usr/local/lib/python3.7/site-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

```
[35]: <AxesSubplot:xlabel='hr', ylabel='cnt'>
```



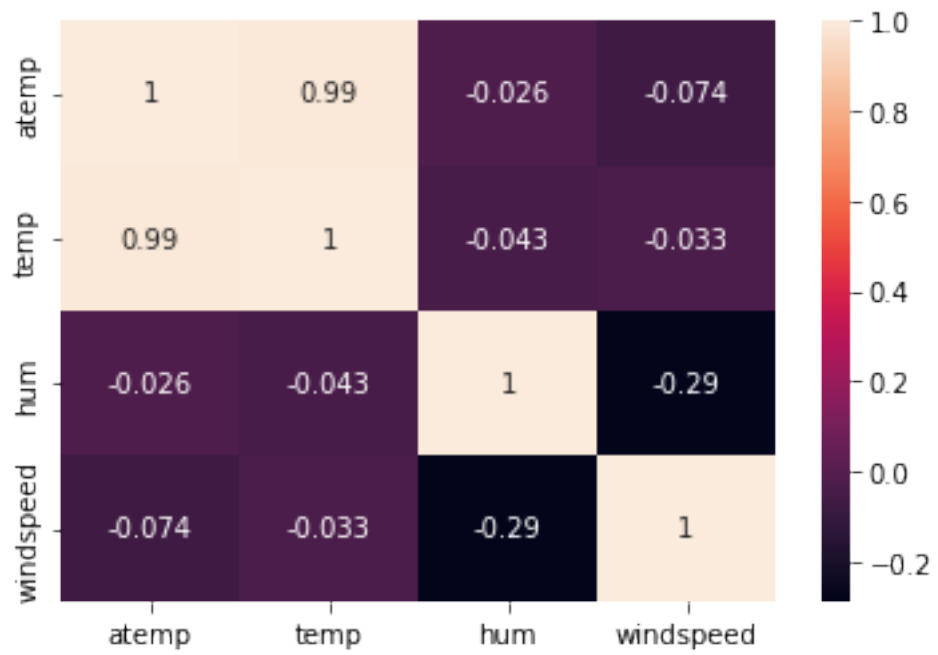
```
[6]: #Make a correlation matrix for variables - atemp, temp, hum, windspeed
#Which variables have the highest correlation?
list1=['atemp','temp','hum','windspeed']
corrs=inp2[list1].corr()
corrs
```

```
[6]:
```

	atemp	temp	hum	windspeed
atemp	1.000000	0.988218	-0.025747	-0.073985
temp	0.988218	1.000000	-0.042603	-0.033209
hum	-0.025747	-0.042603	1.000000	-0.288648
windspeed	-0.073985	-0.033209	-0.288648	1.000000

```
[37]: #create the heat map
sns.heatmap(corrs,annot=True)
```

```
[37]: <AxesSubplot:>
```



- 1 Data pre-processing
- 2 A few key considerations for the pre-processing –
- 3 We seem to have plenty of categorical features.
- 4 Since these categorical features can't be used in the predictive model, we need to convert to a suitable numerical representation.
- 5 Instead of creating dozens of new dummy variables, we will try to club levels of categorical features wherever possible.
- 6 For a feature with high number of categorical levels, we can club the values that are very similar in value for the target variable
- 7 First, create a copy of the dataframe into inp3
- 8 Treating 'mnth' column
- 9 For values 5,6,7,8,9,10 – replace with a single value 5.
- 10 This is because these have very similar values for cnt.
- 11 Get dummies for the updated 6 'mnth' values

```
[39]: inp2.head()
```

```
[39]:
```

	season	yr	mnth	hr	holiday	weekday	workingday	weathersit	temp	\
0	1	0	1	0	0	6	0	1	0.24	
1	1	0	1	1	0	6	0	1	0.22	
2	1	0	1	2	0	6	0	1	0.22	
3	1	0	1	3	0	6	0	1	0.24	
4	1	0	1	4	0	6	0	1	0.24	

	atemp	hum	windspeed	cnt
0	0.2879	0.81	0.0	16
1	0.2727	0.80	0.0	40
2	0.2727	0.80	0.0	32

```
3  0.2879  0.75          0.0   13
4  0.2879  0.75          0.0    1
```

```
[7]: inp3=inp2.copy()
```

```
[8]: #Treating 'mnth' column
#For values 5,6,7,8,9,10 - replace with a single value 5.
#This is because these have very similar values for cnt.
#using isin function
inp3.mnth[inp3.mnth.isin([5,6,7,8,9,10])]=5
```

```
/usr/local/lib/python3.7/site-packages/ipykernel_launcher.py:5:
```

```
SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
"""
```

```
[9]: inp3['mnth'].value_counts()
```

```
[9]: 5      8126
     12     1455
     1      1429
     3      1412
     11     1392
     4      1349
     2      1339
     Name: mnth, dtype: int64
```

```
[10]: #Treating 'hr' column
#Create new mapping: 0-5: 0, 11-15: 11, other values are untouched.
#Again, the bucketing is done in a way that hr values with similar levels of cnt
↪are treated the same.
inp3.hr[inp3.hr.isin([0,1,2,3,4,5])]=0
inp3.hr[inp3.hr.isin([11,12,13,14,15])]=11
```

```
/usr/local/lib/python3.7/site-packages/ipykernel_launcher.py:4:
```

```
SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
after removing the cwd from sys.path.
```

```
/usr/local/lib/python3.7/site-packages/ipykernel_launcher.py:5:
```

```
SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

"""

```
[11]: inp3['hr'].value_counts()
```

```
[11]: 0      4276
      11     3482
      23      728
      22      728
      10      727
      9       727
      21      727
      20      727
      6       725
      7       724
      16      689
      19      671
      8       547
      18      546
      17      478
      Name: hr, dtype: int64
```

```
[12]: #Get dummy columns for season, weathersit, weekday, mnth, hr.
      #We needn't club these further, because as seen from the box plots,
      #the levels seem to have different values for the median cnt.
      list2=['season','weathersit','weekday','mnth','hr']
      inp3=pd.get_dummies(inp3,columns=list2)
```

```
[13]: inp3.head()
```

```
[13]:   yr  holiday  workingday  temp  atemp  hum  windspeed  cnt  season_1 \
0    0         0           0  0.24  0.2879  0.81         0.0   16         1
1    0         0           0  0.22  0.2727  0.80         0.0   40         1
2    0         0           0  0.22  0.2727  0.80         0.0   32         1
3    0         0           0  0.24  0.2879  0.75         0.0   13         1
4    0         0           0  0.24  0.2879  0.75         0.0    1         1

      season_2  ...  hr_10  hr_11  hr_16  hr_17  hr_18  hr_19  hr_20  hr_21 \
0            0  ...     0     0     0     0     0     0     0     0
1            0  ...     0     0     0     0     0     0     0     0
2            0  ...     0     0     0     0     0     0     0     0
3            0  ...     0     0     0     0     0     0     0     0
4            0  ...     0     0     0     0     0     0     0     0

      hr_22  hr_23
```

```

0      0      0
1      0      0
2      0      0
3      0      0
4      0      0

```

[5 rows x 45 columns]

```

[46]: #Train test split - apply 70-30 split
      #call the new dataframes df_train, df_test
      from sklearn.model_selection import train_test_split
      df_train,df_test=train_test_split(inp3,test_size=0.3,random_state=100)

```

```

[47]: df_train.shape

```

```

[47]: (11551, 45)

```

```

[48]: df_test.shape

```

```

[48]: (4951, 45)

```

```

[49]: #Separate X and Y for df_train and df_test.
      #example - you should have X_train, y_train from df_train.
      #y_train should be the cnt column from inp3, X_train should be all other
      ↪ columns.
      y_train=df_train.pop('cnt')
      X_train=df_train

```

```

[50]: y_test=df_test.pop('cnt')
      X_test=df_test

```

```

[51]: X_train

```

```

[51]:      yr  holiday  workingday  temp  atemp  hum  windspeed  season_1  \
9491   1         0             0  0.24  0.2273  0.70    0.2239         1
8763   1         0             1  0.24  0.2576  0.70    0.1045         1
6559   0         0             1  0.50  0.4848  0.82    0.1045         0
2655   0         0             1  0.70  0.6515  0.65    0.2239         0
5646   0         0             0  0.76  0.6818  0.40    0.2985         0
...    ..      ...          ...    ...    ...    ...    ...      ...
17181  1         0             0  0.34  0.3333  0.34    0.1940         1
79     0         0             1  0.22  0.2121  0.51    0.2985         1
12455  1         0             1  0.52  0.5000  0.63    0.0000         0
14761  1         0             1  0.50  0.4848  0.72    0.0896         0
5686   0         0             1  0.66  0.6212  0.50    0.2239         0

      season_2  season_3  ...  hr_10  hr_11  hr_16  hr_17  hr_18  hr_19  \

```

9491	0	0	...	0	0	0	0	0	0
8763	0	0	...	0	0	0	0	0	0
6559	0	0	...	0	0	0	0	0	0
2655	1	0	...	0	1	0	0	0	0
5646	0	1	...	0	0	0	1	0	0
...
17181	0	0	...	0	0	1	0	0	0
79	0	0	...	0	1	0	0	0	0
12455	1	0	...	0	0	0	0	0	0
14761	0	1	...	0	0	0	0	0	0
5686	0	1	...	0	0	0	0	0	0

	hr_20	hr_21	hr_22	hr_23
9491	0	0	0	0
8763	0	0	0	1
6559	0	1	0	0
2655	0	0	0	0
5646	0	0	0	0
...
17181	0	0	0	0
79	0	0	0	0
12455	0	0	0	0
14761	0	0	0	0
5686	0	0	0	0

[11551 rows x 44 columns]

[52]: y_train

```
[52]: 9491      92
      8763      60
      6559     201
      2655     203
      5646     398
      ...
      17181     190
      79       57
      12455     10
      14761     205
      5686     218
      Name: cnt, Length: 11551, dtype: int64
```

[53]: X_test

```
[53]:      yr  holiday  workingday  temp  atemp  hum  windspeed  season_1  \
14790   1         0             1  0.66  0.6212  0.50      0.1642         0
10590   1         0             1  0.50  0.4848  0.94      0.0896         0
```


10635	1	0	0	0.52	0.5000	0.83	0.1642	0
13498	1	0	0	0.60	0.5455	0.88	0.2537	0
9899	1	0	1	0.40	0.4091	0.62	0.2836	1
...
13832	1	0	0	0.88	0.8182	0.42	0.2985	0
1807	0	0	1	0.34	0.3030	0.66	0.3881	0
3598	0	0	1	0.64	0.6212	0.29	0.1642	0
14206	1	0	1	0.60	0.5455	0.88	0.1045	0
15498	1	0	1	0.36	0.3485	0.57	0.1940	0

	season_2	season_3	...	hr_10	hr_11	hr_16	hr_17	hr_18	hr_19	\
14790	0	1	...	0	1	0	0	0	0	
10590	1	0	...	0	0	0	0	0	0	
10635	1	0	...	0	0	0	0	0	0	
13498	0	1	...	0	1	0	0	0	0	
9899	0	0	...	0	1	0	0	0	0	
...	
13832	0	1	...	0	1	0	0	0	0	
1807	1	0	...	0	0	0	0	0	0	
3598	1	0	...	0	0	0	0	0	0	
14206	0	1	...	0	0	0	0	0	0	
15498	0	0	...	0	0	0	0	0	0	

	hr_20	hr_21	hr_22	hr_23
14790	0	0	0	0
10590	0	0	0	0
10635	0	0	0	0
13498	0	0	0	0
9899	0	0	0	0
...
13832	0	0	0	0
1807	0	0	0	0
3598	1	0	0	0
14206	0	0	0	0
15498	0	0	0	1

[4951 rows x 44 columns]

[54]: y_test

[54]: 14790 254
10590 318
10635 28
13498 380
9899 207
...
13832 527

```
1807      13
3598     293
14206      3
15498     100
Name: cnt, Length: 4951, dtype: int64
```

```
[75]: #Model building
      #Use Linear regression as the technique
      #Report the R2 on the train set
      from sklearn.linear_model import LinearRegression
      lin_reg=LinearRegression()
```

```
[76]: #using fit() method for training
      lin_reg.fit(X_train,y_train)
```

```
[76]: LinearRegression()
```

```
[77]: y_pred=lin_reg.predict(X_test)
```

```
[78]: y_pred
```

```
[78]: array([283.1875, 253.5    , 135.625 , ..., 272.125 ,  68.875 , 151.6875])
```

```
[79]: #Reporting r2 for the model
      from sklearn.metrics import r2_score
      print(r2_score(y_pred,y_test))
```

```
0.49236824233484466
```

```
[83]: from sklearn.metrics import r2_score
      print(r2_score(lin_reg.predict(X_train),y_train))
```

```
0.5074631433908097
```

```
[ ]:
```